

Agentic Lead Capture Architecture for the *WhatsApp Diabetes Educator*

Intelligent scoring, persistent user memory, consent-first capture, and structured sales handoff – built on an agentic RAG foundation.

VERSION

v0.3 – Full Architecture

REPLACES

v0.1 Linear Scoring Doc

AUDIENCE

Engineering · Marketing · Sales

STATUS

For Review – Pre-Build

SYSTEM


WhatsApp · Agentic RAG

WHAT THIS DOCUMENT IS

This document defines how the Preventify diabetes educator chatbot should **identify, score, and capture high-quality leads** – without disrupting the clinical trust the chatbot must build with users. It replaces the v0.1 linear scoring model, which did not account for question volume or user history. The architecture described here is **agentic**: the chatbot does not just answer questions – at every turn, it makes decisions about what to do next, guided by what it knows about this specific user across all their conversations. A returning user is never treated as a stranger.

01 The Agentic *Decision Loop* — What Happens at Every Message

Every message a user sends passes through a decision loop before any response is generated. The chatbot is not a simple question-answering machine – it is an agent that **reasons about each message in the context of the full user history** before deciding what to do. This loop runs invisibly on every turn.



WHY THIS MATTERS FOR LEAD QUALITY

The routing decision at step 4 is what separates this from a basic chatbot. The agent never fires a consent prompt mid-answer because a threshold

was crossed – it **waits for a natural pause point** in the conversation where the consent moment will feel earned, not interrupting. This is a judgment call the agent makes, not a mechanical trigger.

02 Question Depth Score — *Classifying Intent*

Every question is assigned a depth score from 1 to 5 by the agent. This is not a keyword match – the agent reads the meaning and context of the question to assign the right level. The same words can score differently depending on what came before in the conversation.

SCORE	INTENT LEVEL	WHAT IT SIGNALS	EXAMPLE QUESTION
1	General Awareness	Definitional, no personal stake. Could be anyone asking out of curiosity.	<i>"What is HbA1c?"</i>
2	Personal Relevance	User applies information to their own situation. Mild but real stake.	<i>"My HbA1c came back at 7.2 — is that okay?"</i>
3	Active Management	User is making decisions about their condition. Real clinical engagement.	<i>"Should I take metformin before or after food?"</i>
4	Complication Concern	User is facing or worried about a complication. High personal urgency.	<i>"My feet go numb at night — is that from diabetes?"</i>
5	Complex / Distressed	Multiple conditions, medication decisions, emotional signals, insulin	<i>"My doctor wants to put me on insulin. I'm scared. What</i>

		management.	<i>should I know?"</i>
--	--	-------------	------------------------

CLASSIFICATION NOTE FOR ENGINEERING

The QDS is assigned by the LLM agent using a classification prompt, not a keyword ruleset. The clinical lead (Dr. Rakesh) must review and validate the classification logic on a sample of 50 real questions before the system goes live. **Misclassification here directly affects lead quality.**

03 Volume Decay — Preventing *Score Inflation*

The core problem with v0.1: a user who asks ten basic awareness questions (QDS 1) accumulates the same score as a user who asks two serious clinical questions (QDS 3-5). This inflates the pipeline with low-quality, unqualified entries. The fix is a **decay multiplier on low-depth questions only**. Questions of depth 3, 4, and 5 always score at full value – they signal genuine intent regardless of volume.

DECAY MULTIPLIER – HOW LOW-DEPTH QUESTIONS ARE WEIGHTED OVER VOLUME

Q1 – Q4 (low depth)	1.0× full value
Q5 – Q7 (low depth)	0.5× half value
Q8+ (low depth)	0.25× minimal
QDS 3, 4, 5 – always	1.0× no decay

With this mechanism, a user who sends ten QDS-1 questions scores approximately 6.5 – safely below the capture threshold of 8. A user who sends two QDS-4 questions scores 8 immediately. The score reflects **intent, not activity**.

94 Persistent User Memory — *Returning Users Are Never Strangers*

A WhatsApp number is a permanent identity. Every time a user messages the chatbot – whether it is their first conversation or their tenth – the agent looks up their full history and continues from where things left off. **There is no concept of a fresh session for a returning user.** Score, condition profile, consent status, and conversation context all persist.

LAYER 1 – IDENTITY

Who is this person

- WhatsApp number (primary key)
- Name (if captured)
- Age (if captured)
- First contact date
- Consent status + timestamp

LAYER 2 – CLINICAL PROFILE

What condition context exists

- Detected diabetes type (T1/T2/GDM/Pre)
- Complications mentioned (kidney, foot, eye)
- Medications referenced
- Highest QDS ever asked
- Clinical escalation history

LAYER 3 – ENGAGEMENT

How engaged are they

- Lifetime score (with decay applied)
- Number of sessions
- Last session date
- Recency weight (see below)
- Lead status (uncaptured / captured / declined)

RECENCY WEIGHTING – HOW RETURNING USER SCORES WORK

A user who asked basic questions six months ago and now returns with complex medication questions is not the same user as their history suggests. The lifetime score applies a **recency weight**: questions from the current session count at full value, questions from the previous 30 days at 0.8×, and older questions at 0.5×. This means a returning user who deepens their engagement will cross the capture threshold faster than their old score alone would suggest – which is correct behaviour.

NEW USER – FIRST MESSAGE

What the agent does

- 1 No memory found for this number – initialise new user profile
- 2 Assign QDS to first question
- 3 Begin building clinical profile from conversation
- 4 Score starts at zero, decay tracking begins
- 5 Respond from RAG corpus as normal

RETURNING USER – ANY SUBSEQUENT MESSAGE

What the agent does

- 1 Number recognised – pull full memory profile
- 2 Load: lifetime score, condition profile, consent status, last session summary
- 3 Apply recency weight to historical score
- 4 Continue conversation with full context – no re-introduction
- 5 If already consented: do not re-prompt. Update lead packet if new info emerges.
- 6 If previously declined: wait until lifetime score exceeds 12 before re-prompting, once only

05 Capture Trigger — The *Dual Condition Gate*

Capture is never triggered by score alone. Two conditions must be true simultaneously. This prevents volume-inflated scores from firing capture on unqualified users, and ensures the captured user has demonstrated real clinical engagement – not just curiosity.

BOTH CONDITIONS MUST BE TRUE – EITHER ALONE IS INSUFFICIENT

CONDITION A – LIFETIME SCORE
THRESHOLD

CONDITION B – DEPTH ANCHOR

The user's **lifetime score reaches 8 or above**, calculated with decay applied to low-depth questions and recency weighting applied to historical questions. Volume of basic questions alone cannot reach this threshold.

AND

The user has asked **at least one question of QDS 3 or higher** at any point in their lifetime conversation history. This confirms the user is actively managing a real condition – not just browsing.

MINIMUM MESSAGE FLOOR

Even if both conditions are met, capture cannot trigger before the user has sent a **minimum of 3 messages** in total (across their lifetime history). This ensures the AI-generated concern summary is meaningful and that the consent moment does not arrive too abruptly.

CLINICAL ESCALATION EXCLUSION

Questions that trigger clinical escalation – hypoglycemic emergencies, chest pain, suicidal ideation, acute symptoms – must **never** be routed into the lead capture flow. These follow the clinical escalation path governed by the Telemedicine Practice Guidelines. The agent must make this routing decision before scoring.

06 The Consent Moment — *Earned, Not Extracted*

When the trigger gate opens, the agent does not immediately interrupt with a consent request. It **waits for a natural conversational pause** – typically after delivering a substantive answer – and then transitions into the consent ask. The consent message must reference something specific from the conversation. A generic "talk to our team" prompt defeats the purpose entirely.

ILLUSTRATIVE CONVERSATION – CONSENT TRIGGER FIRING AFTER A QDS 4 QUESTION

USER

My creatinine is going up and my doctor said my kidneys might be getting affected. I'm on metformin. Should I be worried?

BOT

That is an important concern. When creatinine rises in someone with diabetes, it can signal early kidney involvement – a condition called diabetic nephropathy. Metformin does need to be reviewed in this context, as its use depends on how well the kidneys are functioning. Your doctor is right to be monitoring this closely. The key things to discuss at your next visit are your eGFR number, whether your metformin dose needs adjustment, and whether an SGLT2 inhibitor might be appropriate for kidney protection.

BOT

You've been asking some really thoughtful questions about managing diabetes with kidney involvement. This is exactly the kind of situation where a conversation with a diabetes care advisor – someone who can look at your full picture – tends to be very helpful.

Would it be alright if I shared a brief summary of our conversation with one of our advisors at Preventify? They could follow up with you directly. No obligation – just a conversation if you'd like it.

May I know your name?

USER

Sure, I'm Arjun.

BOT

Thank you Arjun. And your age, so our advisor can give you the most relevant guidance?

Capture triggered here – name + age collected conversationally.
WhatsApp number already available. DPDP consent recorded with timestamp
at first acknowledgement.
AI brief generated immediately from full conversation history and pushed
to sales pipeline.

- 61 The consent message must always reference a **specific topic from the conversation**. The agent generates this contextually – it is never a templated string.
-
- 62 Only name and age are collected in-chat. WhatsApp number is already captured. **No email, address, or financial information** at this stage.
-
- 63 If the user declines, the chatbot continues normally. **No re-prompt within the same session**. For returning users who have previously declined, see Section 04.
-
- 64 Consent language must be DPDP Act 2023 compliant. The user must be told **explicitly what their data will be used for** and by whom. Consent timestamp is stored in the user memory profile.
-
- 65 The agent must **never ask for consent mid-answer**. The full clinical answer must be delivered first. The consent transition only happens at a natural pause.

07 The AI Brief — *What Gets Handed to the Sales Team*

The AI brief is generated by the agent at the moment consent is given. It draws on the **full conversation history across all sessions** – not just the current conversation. The sales advisor should be able to open this brief and make a confident, personalised call without reading a transcript. The brief is a warm handoff note, not a data dump.

01 · IDENTITY

Name · Number · Age

Collected at consent. WhatsApp number is the persistent identifier – every other field is built on top of it.

02 · CONDITION TYPE

Detected Diabetes Profile

T1 / T2 / GDM / Prediabetes / Undiagnosed – inferred by the agent from conversation context. Tells the advisor which

03 · CONCERN SUMMARY

2–3 Line AI-Generated Brief

e.g. "User managing T2DM with early kidney involvement. Asking about metformin safety and eGFR thresholds. Has been

	Preventify service to discuss.	tracking the condition actively for some time."
--	--------------------------------	---

04 - ENGAGEMENT SCORE Lifetime Score at Capture The score that triggered capture. Tells the advisor how warm this lead is – a score of 14 is a different conversation opener than a score of 8.	05 - PEAK CONCERN Highest QDS Asked (Lifetime) The most serious topic raised across all sessions. Gives the advisor their first conversation anchor – the thing the user cares about most.	06 - CAPTURE TIMESTAMP Date + Time (IST) When consent was given. Used for follow-up SLA and pipeline prioritisation. Older uncalled leads should be flagged.
---	--	--

BRIEF QUALITY REQUIREMENT The concern summary must be generated from actual conversation content by the LLM – not templated, not rule-based. A templated brief defeats the purpose of personalised outreach. The brief should read like a colleague's handoff note, not a form field. Engineering must treat brief quality as a first-class output alongside answer quality.

08 Sales Pipeline Structure — *Built From Scratch*

Preventify does not currently have a structured CRM. This is an opportunity to define a clean, minimal pipeline from the ground up – one that matches the nature of leads coming from a health chatbot. The recommended structure is four stages, reflecting the natural progression of a health-service sales conversation.

STAGE 1 New Lead Consent given. AI brief	STAGE 2 Contacted Advisor has called or	STAGE 3 Qualified User has expressed	STAGE 4 Converted / Closed User has
--	---	--	---

generated. Not yet contacted by the sales team. Triggered by: consent capture in chatbot	messed. Conversation has begun. Outcome not yet known. Triggered by: first advisor outreach logged	genuine interest in a Preventify service. Next step agreed. Triggered by: advisor marks as qualified	enrolled in or purchased a Preventify service. Or marked as not converting. Triggered by: enrolment or advisor closure
---	---	---	---

RECOMMENDED TOOLING

For an unstructured pipeline starting from zero, a lightweight CRM like **Zoho CRM Free** or **HubSpot Starter** is recommended over a spreadsheet. The AI brief fields map cleanly to custom CRM properties. The chatbot pushes new leads via a webhook at the moment consent is captured – no manual entry required. The sales team views and updates leads entirely within the CRM.

09 Next Steps — *By Team*

ENGINEERING	CLINICAL + MARKETING	SALES + OPERATIONS
→ Define user memory schema (identity, clinical profile, engagement layers)	→ Review and sign off on QDS classification examples (Dr. Rakesh)	→ Select and set up CRM (Zoho or HubSpot recommended)
→ Build QDS classification prompt and validate with clinical lead on 50 sample questions	→ Review consent message language for tone and compliance	→ Define SLA for first contact after lead capture (recommended: within 24 hours)
→ Implement decay multiplier and recency weighting in scoring logic	→ Define the Preventify service that each condition type (T1/T2/GDM)	→ Train advisors on reading and using the AI brief format
		→ Define pipeline stage criteria and who

→ Build dual condition gate with minimum message floor

→ Build AI brief generation prompt and test output quality

→ Set up webhook from chatbot to CRM pipeline on consent trigger

→ Implement DPDP-compliant consent logging with timestamp

maps to

→ Agree on re-prompt threshold for users who previously declined

→ Review 20 sample AI briefs before go-live for quality and tone

moves leads between stages

→ Set up monthly review of lead quality vs. conversion to tune scoring thresholds

Preventify – Diabetes Educator Intelligence

Layer v0.3

Companion documents: Guideline Corpus Spec ·

RAG Architecture Doc

For: Engineering + Clinical Lead + Sales
review before build

V0.3 · WORKING DRAFT · OPEN
FOR REVIEW